Inverse Estimation with Linear Mixed-Effects Models

Brandon M. Greenwell
Infoscitex, a DCS company
and
Christine M. Schubert Kabban
Air Force Institute of Technology

October 23, 2016

Abstract

Keep it to 200 words or less.

1 Introduction

This paper concerns inverse estimation with linear mixed-effects models (LMMs). Consider an ordinary regression model $\mathcal{Y}_i = f(x_i; \boldsymbol{\beta}) + \epsilon_i$ $(i = 1, \ldots, n)$, where f is a known expectation function (called a *calibration curve*) that is monotonic over the range of interest and $\epsilon_i \overset{iid}{\sim} \mathcal{N}(0, \sigma^2)$. A common problem in regression is to predict a future response $\mathcal{Y}_0$ (or estimate the mean response $f_0 = \mathrm{E}[\mathcal{Y}_0]$) for a known value of the explanatory variable x_0. Often, however, there is a need to do the reverse; that is, given an observed value of the response $\mathcal{Y} = y_0$ (or a specified value of the mean response), infer the unknown value of the explanatory variable x_0. This is known as the *calibration problem*, though we refer to it more generally as inverse estimation. A thorough overview of the calibration problem is given in Osborne [1991]. Oman [1998] considers the case of a random intercept and slope model.

In Section 3, we extend the classical inverse estimate. Section 4 and Section 5 discuss approximate interval estimates by extending the classical Wald and inversion methods, respectively. Section 6 introduces a fully parametric bootstrap algorithm for making inference on x_0.

2 Linear mixed-effects models

Linear regression models with random coefficients can be represented in many different (but equivalent) forms. One of the most common forms, attributed to

Laird and Ware [1982], is

$$\boldsymbol{\mathcal{Y}}_i = \boldsymbol{X}_i \boldsymbol{\beta} + \boldsymbol{Z}_i \boldsymbol{b}_i + \boldsymbol{\epsilon}_i, \quad i = 1, \ldots, m, \qquad (1)$$

where

- $\boldsymbol{\mathcal{Y}}_i$ is an $n_i \times 1$ response vector for the i-th subject/cluster/group;

- $\boldsymbol{X}_i$ is an $n_i \times p$ design matrix for the fixed-effects;

- $\boldsymbol{Z}_i$ is an $n_i \times q$ design matrix for the random-effects;

- $\boldsymbol{\beta}$ is a $p \times 1$ vector of fixed-effects coefficients;

- $\boldsymbol{b}_i$ is a $q \times 1$ vector of random-effects coefficients with mean zero and variance-covariance matrix $\boldsymbol{D}$;

- $\boldsymbol{D}$ is a $q \times q$ variance-covariance matrix for the random-effects;

- $\boldsymbol{\epsilon}_i$ is an $n_i \times 1$ vector of random errors with mean zero and variance-covariance matrix $\sigma^2 \boldsymbol{I}$.

Equation (1) is known as a linear mixed-effects model (LMM). The random-effects and errors are often assumed to follow a normal distribution. By stacking the data, the (normal) LMM can be written concisely as

$$\boldsymbol{\mathcal{Y}} = \boldsymbol{X}\boldsymbol{\beta} + \boldsymbol{Z}\boldsymbol{b} + \boldsymbol{\epsilon}, \quad \begin{bmatrix} \boldsymbol{b} \\ \boldsymbol{\epsilon} \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} \boldsymbol{0} \\ \boldsymbol{0} \end{bmatrix}, \begin{bmatrix} \boldsymbol{D} & \boldsymbol{0} \\ \boldsymbol{0} & \sigma^2 \boldsymbol{I} \end{bmatrix} \right), \qquad (2)$$

where $\boldsymbol{\mathcal{Y}} = \mathrm{col}\{\boldsymbol{\mathcal{Y}}_i\}$, $\boldsymbol{X} = \mathrm{col}\{\boldsymbol{X}_i\}$, $\boldsymbol{Z} = \mathrm{diag}\{\boldsymbol{Z}_i\}$, $\boldsymbol{b} = \mathrm{col}\{\boldsymbol{b}_i\}$, and $\boldsymbol{\epsilon} = \mathrm{col}\{\boldsymbol{\epsilon}_i\}$ for $i = 1, \ldots, m$. Since $\mathrm{COV}\,[\boldsymbol{b}, \boldsymbol{\epsilon}] = \boldsymbol{0}$, it is assumed that the random vectors $\{\boldsymbol{b}_i, \boldsymbol{\epsilon}_i\}_{i=1}^m$ are mutually independent.

The additional term $\boldsymbol{Z}\boldsymbol{b}$ in the model imposes a specific variance-covariance structure on the response vector $\boldsymbol{\mathcal{Y}}$:

$$\boldsymbol{\mathcal{Y}} \sim \mathcal{N}\left(\boldsymbol{X}\boldsymbol{\beta}, \boldsymbol{V}\right), \quad \boldsymbol{V} = \boldsymbol{Z}\boldsymbol{D}\boldsymbol{Z}^\top + \sigma^2 \boldsymbol{I}.$$

Thus, the fixed-effects determine the mean of $\boldsymbol{\mathcal{Y}}$, while the random-effects govern the variance-covariance structure of $\boldsymbol{\mathcal{Y}}$. Different random-effects structures impose different variance-covariance structures on the response resulting in a highly flexible framework for modelling *grouped data*.

The random-effects variance-covariance matrix $\boldsymbol{D}$ has at most $q(q+1)/2$ unique elements which we represent by the vector $\boldsymbol{\theta}$. There are a number of methods available for estimating $(\boldsymbol{\beta}, \sigma^2, \boldsymbol{\theta})$; see, for example, McCulloch et al. [2008, chap. 6.] and Demidenko [2013, chap. 2]. Most commonly, the fixed-effects $\boldsymbol{\beta}$ are estimated via the method of maximum likelihood (ML), while the variance components $(\sigma^2, \boldsymbol{\theta})$ are estimated via restricted maximum likelihood (REML). The ML estimator of $\boldsymbol{\beta}$, given by

$$\widehat{\boldsymbol{\beta}} = \left(\boldsymbol{X}^\top \widehat{\boldsymbol{V}}^{-1} \boldsymbol{X} \right) \boldsymbol{X}^\top \boldsymbol{\mathcal{Y}},$$

depends on the estimated variance components through $\widehat{V}$ which makes it difficult to capture the variability of $\widehat{\boldsymbol{\beta}}$ in small sample sizes (see McCulloch et al. [2008, pp. 165-167]). The usual practice is to ignore the variability of the estimated variance components when making inference about the fixed-effects; that is, treat $\widehat{V}$ as the true (fixed) value of V. Modern computational procedures such as the parametric bootstrap and Markov chain Monte Carlo (MCMC) methods are two ways account for the variability of the estimated variance components.

2.1 Bladder volume data

For illustration, let us consider the bladder volume data which can be found in Brown [1993, pg. 7] and Oman [1998]. In Brown's words:

> "A series of 23 women patients attending a urodynamic clinic were recruited for the study. After successful voiding of the bladder, sterile water was introduced in additions of 1, 1.5, and then 2.5 cl increments up to a final cumulative total of 17.5 cl. At each volume a measure of height (H) in mm and depth (D) in mm of largest ultrasound bladder images were taken. The product H × D was taken as a measure of liquid volume."

We took Brown's suggestion and transformed the data so that the relationship is approximately linear. Spaghettiplots of both the original and transformed data are displayed in Figure 1.


Figure 1: Spaghettiplot of bladder volume data (lines connect measurements belonging to the same subject). *Left*: Original data. *Right*: Transformed data.

A random intercept and slope model with uncorrelated random-effects fits

Parameter	Value	Std.Error	DF	t-value	p-value
intercept	-53.832	11.603	142	-4.639	< 0.001
volume	69.095	3.114	142	22.191	< 0.001

Table 1: Output for fixed-effects portion of LMM fit to the transformed bladder volume data.

the transformed data (right side of Figure 1) well:

$$\mathtt{HD}_{ij}^{3/2} = (\beta_0 + b_{0i}) + (\beta_1 + b_{1i})\, \mathtt{volume}_{ij} + \epsilon_{ij}$$
$$b_{ki} \sim \mathcal{N}\left(0, \theta_k^2\right), \quad k = 0, 1,$$
$$\epsilon_{ij} \sim \mathcal{N}\left(0, \sigma^2\right),$$

where $\mathrm{COV}\,[b_{0i}, b_{1i}] = 0$. Table 1 displays the fixed-effects results from applying this model to the transformed bladder volume data.

For an in-depth treatment on fitting LMMs using the `nlme` software, see Pinheiro and Bates [2000].

3 Point estimation

The standard methods of calibration, (i.e., the Wald-based and inversion confidence intervals) are easily extended to the case of random coefficients. For convenience, let us write the linear random coefficient model as

$$\mathcal{Y}_{ij} = f\left(x_{ij}; \boldsymbol{\beta}\right) + R\left(x_{ij}; \boldsymbol{b}_i\right) + \epsilon_{ij},$$

where $f(\cdot)$ and $R(\cdot)$ are linear in $\boldsymbol{\beta}$ and $\boldsymbol{b}_i$, respectively. For instance, the model for the transformed bladder data has $f\left(\mathtt{volume}_{ij}; \boldsymbol{\beta}\right) = \beta_0 + \beta_1 \mathtt{volume}_{ij}$ and $R\left(\mathtt{volume}_{ij}; \boldsymbol{b}_i\right) = b_{0i} + b_{1i}\mathtt{volume}_{ij}$ with $\mathrm{E}\left[R\left(\mathtt{volume}_{ij}; \boldsymbol{b}_i\right)\right] = 0$ and $\mathrm{VAR}\left[R\left(\mathtt{volume}_{ij}; \boldsymbol{b}_i\right)\right] = \theta_0^2 + \mathtt{volume}_{ij}^2 \theta_1^2$.

Assume that, after the data are collected and a model is fitted, we obtain a new observation, denoted $\mathcal{Y}_0$, from the same population under study for which the value of the explanatory variable x_0 is unknown. We assume that the new observation belongs to a group not included in our analysis. Estimating x_0 is rather straightforward. By assumption, the new observation $\mathcal{Y}_0$ is distributed as a $\mathcal{N}\left\{f\left(x_0; \boldsymbol{\beta}\right), \sigma_0^2\right\}$ random variable with $\sigma_0^2 = \mathrm{VAR}\left[R\left(x_0; \boldsymbol{b}_0\right)\right] + \sigma^2$. A natural estimator for x_0 is then

$$\widehat{x}_0 = f^{-1}\left(\mathcal{Y}_0; \widehat{\boldsymbol{\beta}}\right), \tag{3}$$

where $\widehat{\boldsymbol{\beta}}$ is the ML estimator of $\boldsymbol{\beta}$. We shall refer to Equation (3) as the classical estimator. Note that the point estimate $\widehat{x}_0$ does not involve any of the random-effects; the random-effects only contribute to the variance-covariance structure of the response.

4 Wald interval

An approximate $100(1-\alpha)\%$ Wald-type confidence interval for x_0 has the simple form

$$CI_{wald}\left(x_0\right) = \left(\widehat{x}_0 - \text{SE}\left[\widehat{x}_0\right]\Phi\left(\alpha/2\right), \widehat{x}_0 - \text{SE}\left[\widehat{x}_0\right]\Phi\left(1-\alpha/2\right)\right). \qquad (4)$$

There is no "textbook" formula for the standard error of $\widehat{x}_0$—not even in the case of the simple linear regression model. Instead, an estimate of the standard error is obtained using a first-order Taylor series approximation, or better yet, a bootstrap approximation.

The Taylor series approximation relies on the variance-covariance matrix of $\left(\mathcal{Y}_0, \widehat{\boldsymbol{\beta}}\right)$, namely,

$$\Sigma = \begin{bmatrix} \text{VAR}\left[\mathcal{Y}_0\right] & \mathbf{0} \\ \mathbf{0} & \text{VAR}\left[\widehat{\boldsymbol{\beta}}\right] \end{bmatrix} = \begin{bmatrix} \sigma_0^2 & \mathbf{0} \\ \mathbf{0} & \left(\boldsymbol{X}^\top\boldsymbol{V}^{-1}\boldsymbol{X}\right)^{-1} \end{bmatrix}.$$

Since $\mathcal{Y}_0$ is independent of $\boldsymbol{\mathcal{Y}}$, it is also independent of $\widehat{\boldsymbol{\beta}}$, hence the diagonal structure of Σ. Recall that our point estimate has the form $x = f^{-1}\left(y; \boldsymbol{\beta}\right)$. Let $f_1^{-1}\left(y; \boldsymbol{\beta}\right)$ and $f_2^{-1}\left(y; \boldsymbol{\beta}\right)$ denote the partial derivatives of f^{-1} with respect to the parameters y and $\boldsymbol{\beta}$, respectively. Our point estimator is given by $f^{-1}\left(\mathcal{Y}_0; \widehat{\boldsymbol{\beta}}\right)$, where $\mathcal{Y}_0$ is a new observation and $\widehat{\boldsymbol{\beta}}$ is the ML estimator of $\boldsymbol{\beta}$. A first-order Taylor-series approximation for the variance of $\widehat{x}_0$ is given by

$$\begin{aligned} \text{VAR}\left[\widehat{x}_0\right] = {} & \left[f_1^{-1}\left(\mathcal{Y}_0; \widehat{\boldsymbol{\beta}}\right)\right]^2 \sigma_0^2 \\ & + \left[f_2^{-1}\left(\mathcal{Y}_0; \widehat{\boldsymbol{\beta}}\right)\right]^\top \left(\boldsymbol{X}^\top\boldsymbol{V}^{-1}\boldsymbol{X}\right)^{-1} \left[f_2^{-1}\left(\mathcal{Y}_0; \widehat{\boldsymbol{\beta}}\right)\right]. \end{aligned} \qquad (5)$$

To obtain $\text{SE}\left[\widehat{x}_0\right] = \left\{\widehat{\text{VAR}}\left[\widehat{x}_0\right]\right\}^{1/2}$, we simply replace σ_0^2 and $\boldsymbol{V}$ with their respective estimates $\widehat{\sigma}_0^2$ and $\widehat{\boldsymbol{V}}$.

The Wald-based interval is simple to compute as long as we have an estimate for the standard error. As we will discuss in Section A, the R package investr [Greenwell, 2013] can be used to obtain the Wald-based interval (4) using a Taylor series approximation of the standard error. If a closed-form formula is available for $\widehat{x}_0$, then the deltaMethod function from the car package [Fox and Weisberg, 2011] can also be used to obtain the Taylor series approximation of the standard error. Alternatively, one can use the parametric bootstrap instead of relying on a Taylor series approximation; see Section 6. The bootstrap estimate may be more accurate in smaller sample sizes because, unlike the Taylor approximation estimate, it takes into account the variability of the estimated variance components. The bootMer function in the lme4 package [Bates et al., 2014] can be used for model-based parametric bootstrapping in mixed-effects models.

5   Inversion interval

In the case of the simple linear regression model with constant variance, an exact $100(1-\alpha)\%$ confidence interval for x_0 can be derived [Graybill, 1976]. This can be generalized to an approximate method in the case of polynomial or nonlinear regression models with independent observations and constant variance (see Seber and Wild [2003] and Huet [2004]). In a similar fashion, we can generalize the same results to an approximate method for linear mixed-effects models.

Let $\widehat{f}_0 = f\left(x_0; \widehat{\boldsymbol{\beta}}\right)$ be the predicted mean at $x = x_0$. A prediction interval for $\mathcal{Y}_0$ at x_0 with asymptotic coverage probability $100(1-\alpha)\%$ is

$$\mathcal{I}_\infty\left(x_0\right) = \widehat{f}_0 \pm z_{1-\alpha/2}\left\{\widehat{\mathrm{VAR}}\left[\mathcal{Y}_0 - \widehat{f}_0\right]\right\}^{1/2}. \tag{6}$$

If instead, $\mathcal{Y}_0$ is observed to be y_0 and x_0 is unknown, then an asymptotic $100(1-\alpha)\%$ confidence interval for the unknown x_0 can be obtained by inverting (6):

$$CI_{inv}\left(x_0\right) = \left\{x : z_{\alpha/2} \leq \frac{\mathcal{Y}_0 - f\left(x; \widehat{\boldsymbol{\beta}}\right)}{\left\{\widehat{\mathrm{VAR}}\left[\mathcal{Y}_0 - f\left(x; \widehat{\boldsymbol{\beta}}\right)\right]\right\}^{1/2}} \leq z_{1-\alpha/2}\right\}. \tag{7}$$

This is known as the *inversion interval* and typically cannot be written in closed-form; therefore, numerical techniques are required to find the lower and upper bounds. Further, note that $CI_{inv}\left(x_0\right)$ is not symmetric about $\widehat{x}_0$ and will not necessarily result in a single finite interval.

Fortunately, the inversion interval (7) can be computed automatically using the `investr` package. However, like the Wald-based interval (4), the inversion interval ignores the variability of the estimated variance components and will likely perform poorly in small sample sizes. An alternative approach involving the parametric bootstrap will be discussed in Section 6.

Finally, the inversion interval uses a normal approximation. While it is likely that a t-distribution may be more accurate, it is difficult to find the appropriate degrees of freedom. Oman [1998], suggests a t-distribution with $N-1$ degrees of freedom (N being the total sample size).

6   Parametric bootstrap

The bootstrap [Efron, 1979] is a general-purpose computer-based method for assessing accuracy of estimators and forming confidence intervals for parameters. Jones and Rocke [1999] proposed a nonparametric bootstrap algorithm for controlled calibration with independent observations. However, since our application involves random coefficients (i.e., dependent observations), the nonparametric bootstrap does not easily apply, and instead, we adopt a "fully parametric" approach. In a parametric bootstrap, bootstrap samples are generated from a fitted parametric model rather than sampling with replacement

directly from the data. Fortunately, parametric bootstrap confidence intervals are usually more accurate than nonparametric ones, however, by sampling from a fitted parametric family, we are implicitly assuming that we have the "correct model".

Let $\widehat{\sigma}_0^2$ be an estimate of the variance of the new observation $\mathcal{Y}_0$. An algorithm for bootstrapping $\widehat{x}_0$ in an LMM is given in Figure 2. Note that step 5. is crucial for calibration problems because we need to treat y_0 as a random quantity in the bootstrap simulation, otherwise the variability of $\widehat{x}_0$ will be underestimated; see, for example, Jones and Rocke [1999] and Greenwell and Kabban [2014].

1. Fit a mixed model (2) to the data and obtain estimates $\widehat{\boldsymbol{\beta}}$, $\widehat{\boldsymbol{D}}$, and $\widehat{\sigma}^2$.

2. Define $\boldsymbol{y}^\star = \boldsymbol{X}\widehat{\boldsymbol{\beta}} + \boldsymbol{Z}\boldsymbol{b}^\star + \boldsymbol{\epsilon}^\star$, where $\boldsymbol{b}^\star \sim \mathcal{N}_q\left(\boldsymbol{0}, \widehat{\boldsymbol{D}}\right)$ and $\boldsymbol{\epsilon}^\star \sim \mathcal{N}_N\left(\boldsymbol{0}, \widehat{\sigma}_\epsilon^2 \boldsymbol{I}\right)$;

3. Update the original model using $\boldsymbol{y}^\star$ as the response vector to obtain $\widehat{\boldsymbol{\beta}}^\star$ and $\widehat{\sigma}_0^{2\star}$;

4. Generate $y_0^\star \sim \mathcal{N}\left(y_0, \widehat{\sigma}_0^{2\star}\right)$;

5. Define $\widehat{x}_0^\star = f^{-1}\left(y_0^\star; \widehat{\boldsymbol{\beta}}^\star\right)$;

6. Repeat steps (2)-(5) R times.

Figure 2: Parametric bootstrap algorithm for linear calibration with random coefficients.

There are three main bootstrap confidence interval procedures: The percentile methods introduced in Efron [1979], the studentized bootstrap t method introduced in Efron [1982], and the double bootstrap method [Hall, 1986]. For a good overview of all these confidence interval procedures, see Davison and Hinkley [1997, chap. 5] and Boos and Stefanski [2013, chap. 11]. In the next section, we discuss how to use this algorithm to adjust the previously discussed inversion interval.

6.1 A bootstrap adjusted inversion interval for x_0

Huet [2004] suggests a bootstrap modification of the usual inversion interval in nonlinear regression models with dependent data. In a similar fashion, we could use the parametric bootstrap to adjust the approximate inversion interval given in Equation (7). The inversion interval assumes that the *predictive pivot*

$$Q_I = \frac{\mathcal{Y}_0 - f\left(x; \widehat{\boldsymbol{\beta}}\right)}{\left\{\widehat{\text{VAR}}\left[\mathcal{Y}_0 - f\left(x; \widehat{\boldsymbol{\beta}}\right)\right]\right\}^{1/2}} \sim \mathcal{N}(0, 1).$$

A bootstrap modified inversion interval would then use the bootstrap distribution of

$$Q_I^\star = \frac{\mathcal{Y}_0^\star - f\left(\widehat{x}_0; \widehat{\boldsymbol{\beta}}^\star\right)}{\left\{\widehat{\mathrm{VAR}}\left[\mathcal{Y}_0 - f\left(\widehat{x}_0; \widehat{\boldsymbol{\beta}}^\star\right)\right]\right\}^{1/2}},$$

to estimate the true distribution of Q_I. If $\widehat{F}_{Q_I}$ is the empirical distribution function for a sample of R bootstrap replicates of Q_I, then the modified inversion interval for x_0 is given by

$$CI_{inv}^\star(x_0) = \left\{x : \widehat{F}_{Q_I}(\alpha/2) \leq Q_I \leq \widehat{F}_{Q_I}(1 - \alpha/2)\right\}.$$

All of the approximate 95% confidence intervals we computed for the true volume of fluid are summarized in Table 2 below.

Method	Estimate	SE	95% Bounds	Length
$CI_{wald}(x_0)$	8.016	1.954	(4.185, 11.846)	7.66
$CI_{inv}(x_0)$	8.016	—	(4.228, 11.919)	7.691
$CI_{percentile}^\star(x_0)$	8.016	1.974	(4.05, 11.913)	7.863
$CI_{wald}^\star(x_0)$	8.016	1.974	(4.252, 11.947)	7.695
$CI_{inv}^\star(x_0)$	8.016	1.974	(4.278, 12.019)	7.741

Table 2: Summary of results for the bladder volume example. A $\star$ symbol indicates a parametric bootstrap-based confidence interval.

Notice that all of the bootstrap-based confidence intervals for x_0 (indicated by a $\star$) are slightly wider than those based on large sample normal theory results. This is likely due to the fact that the large sample intervals do not take into account the variability of the estimated variance components.

6.2 Monte Carlo Study

To assess the empirical performance of these confidence intervals, we carried out a small Monte Carlo study[1]. The results are reported in Table 3 and indicate that the Wald-based confidence interval (4) and inversion confidence interval (7) have asymptotic coverage probability close to $100(1-\alpha)\%$. This experiment also highlighted the fact that it is the number of subject m, not the sample size per subject n, that is more important for good asymptotic coverage. The code used for the simulation is available upon request.

We consider the values 5, 10, 30, 50, and 100 for both the number of subjects m and and the number of observations per subject n. For each combination of

---

[1]The simulation described in this section was conducted in R using packages plyr [Wickham, 2011], nlme, and lme4 [Bates et al., 2014].

sample sizes, we generated 1,000 data sets from a random intercept and slope model with parameters given by those listed in `summary(transformed.lme)`. In other words, the fixed-effects were $\boldsymbol{\beta} = (-53.83164, 69.09491)^\top$, the standard deviations for the (uncorrelated) random intercept and slope were 39.62499, and 14.28841, respectively. The residual standard deviation was $\sigma = 53.71511$. We chose $f(x_0; \boldsymbol{\beta}) = 500$ so that the true unknown is $x_0 = 8.0155$. The standard deviation of the coverage estimates is approximately $\sqrt{0.95(1 - 0.95)/1000} = 0.001$. A trellis plot of the results is given in Figure 3. The coverage estimates are plotted against the number of subjects m and paneled by number of observations per subject n. The results indicate that the number of subject, m, is the real driver in terms of asymptotic coverage probability. It seems that $m \geq 30$ with $n \geq 5$ is sufficient for achieving close to the stated $1 - \alpha$ coverage probability of.

m	Method	$n = 5$	$n = 10$	$n = 30$	$n = 50$	$n = 100$
5	Wald	0.89	0.89	0.89	0.87	0.89
	Inversion	0.89	0.90	0.89	0.88	0.90
10	Wald	0.92	0.92	0.94	0.93	0.93
	Inversion	0.92	0.92	0.94	0.94	0.93
30	Wald	0.95	0.95	0.95	0.94	0.95
	Inversion	0.94	0.94	0.94	0.94	0.95
50	Wald	0.95	0.96	0.95	0.94	0.94
	Inversion	0.94	0.95	0.95	0.94	0.94
100	Wald	0.95	0.95	0.95	0.94	0.94
	Inversion	0.95	0.95	0.95	0.94	0.95

Table 3: Coverage probability of 95% confidence intervals for simulated bladder data.

7 Conclusion

We have discussed a number of confidence interval procedures for statistical calibration in linear models with random coefficients with a single level of grouping. We have described two `R` packages for implementing these procedures: `investr` and `lme4`. The `investr` package can be used for obtaining the asymptotic confidence intervals (i.e., the Wald-based and inversion confidence intervals). We also showed how the `lme4` package can be used to obtain calibration intervals based on a parametric bootstrap using the recently added `bootMer` function. Future work will likely extend the methods discussed in this paper to more complicated cases such as nonlinear mixed-effects models and multi-level hierarchical models (i.e., more than one grouping variable).

Figure 3: Coverage probability of 95% confidence intervals for simulated bladder data. The coverage estimates based on the inversion method are colored blue.

References

Douglas Bates, Martin Maechler, Ben Bolker, and Steven Walker. *lme4: Linear mixed-effects models using Eigen and S4*, 2014. URL http://CRAN.R-project.org/package=lme4. R package version 1.1-6.

D.D. Boos and L.A. Stefanski. *Essential Statistical Inference: Theory and Methods*. Springer Texts in Statistics. Springer London, Limited, 2013.

P.J. Brown. *Measurement, Regression, and Calibration*. Oxford science publications. Clarendon Press, 1993.

A. C. Davison and D. V. Hinkley. *Bootstrap Methods and their Applications*. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, Cambridge, 1997.

E. Demidenko. *Mixed Models: Theory and Applications with R*. Wiley, 2013.

B. Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

B. Efron. *The Jackknife, the Bootstrap, and Other Resampling Plans*. Society for Industrial and Applied Mathematics, 1982.

John Fox and Sanford Weisberg. *An R Companion to Applied Regression*. Sage, Thousand Oaks CA, second edition, 2011. URL http://socserv.socsci.mcmaster.ca/jfox/Books/Companion.

F. A. Graybill. *Theory and Application of the Linear Model*. Duxbury classic series. Duxbury, Pacific Grove, CA, 1976.

Brandon M. Greenwell. *investr: Functions for Inverse Estimation with Linear and Nonlinear Regression Models in R.*, 2013. URL http://CRAN.R-project.org/package=investr. R package version 1.1.0.

Brandon M. Greenwell and Christine M. Schubert Kabban. investr: An r package for inverse estimation. *The R Journal*, 6(1):??–??, 2014.

Peter Hall. On the bootstrap and confidence intervals. *The Annals of Statistics*, 14(4):1431–1452, 1986.

S. Huet. *Statistical Tools for Nonlinear Regression: A Practical Guide with S-PLUS and R Examples*. Springer Series in Statistics. Springer, 2004.

Geoffrey Jones and David M. Rocke. Bootstrapping in controlled calibration experiments. *Technometrics*, 41(3):224–233, 1999.

Nan M. Laird and James H. Ware. Random-effects models for longitudinal data. *Biometrics*, 38(4):963–974, 1982.

C. E. McCulloch, S. R. Searle, and J. M. Neuhaus. *Generalized, Linear, and Mixed Models*. Wiley, 2008.

Samuel D. Oman. Calibration with random slopes. *Biometrika*, 85(2):439–449, 1998.

C. Osborne. Calibration: A review. *International Statistical Review*, 59(3):309–336, 1991.

J.C. Pinheiro and D.M. Bates. *Mixed-Effects Models in S and S-PLUS*. Statistics and Computing. Springer, 2000.

G.A.F. Seber and C.J. Wild. *Nonlinear Regression*. Wiley Series in Probability and Statistics. Wiley, 2003.

Hadley Wickham. The split-apply-combine strategy for data analysis. *Journal of Statistical Software*, 40(1):1–29, 2011. URL http://www.jstatsoft.org/v40/i01/. R package version 1.8.1.

A Inverse estimation in `R`

Here, we discuss how to implement the previous procedures in the `R` programming languages. We discuss two main packages: `investr` and `lme4`. The `investr` package can be used for obtaining the Wald-based and inversion intervals. This functionality is demonstrated over the next two sections. The `lme4` package is a popular package for fitting linear, generalized linear, and nonlinear mixed models. Recently, however, the `lme4` package creators have added functionality for model-based parametric bootstrapping. In Section A.2, we demonstrate the potential of this new functionality by applying our parametric bootstrap algorithm to the bladder volume data discussed earlier.

A.1 The `investr` package

The `R` package `investr` facilitates calibration/inverse estimation with linear and nonlinear regression models. The main function, `invest`, can be used for inverse estimation of x_0 given an observed response y_0. More recently, the package has been updated to also handle objects of class `lme` from the `nlme` package. Current functionality includes both the Wald-based and inversion methods outlined in Sections 4-5. The code for the package is hosted on GitHub at https://github.com/bgreenwell/investr, but the latest stable release can be found on CRAN at https://CRAN.R-project.org/package=investr.

Returning to the bladder volume example, suppose we obtained an ultrasound measurement from a new patient for which $\mathtt{HD}^{3/2} = 500$ (that's roughly 63 on the original scale). What is the true volume of fluid (x_0) in the patients bladder? We can estimate the true volume and form an approximate 95% confidence interval using the methods discussed previously. The point estimate is simply given by

$$\widehat{x}_0 = \frac{500 + 53.83164}{69.09491} = 8.0155 \text{ (cl)}.$$

This estimate can be obtained in `R` as follows:

```
library(investr)
invest(transformed.lme, y0 = 500, interval = "none")

## [1] 8.02
```

`invest` relies on the root finding function `uniroot` from the `stats` package to solve the equation $f\left(x; \widehat{\boldsymbol{\beta}}\right) - y_0 = 0$ numerically for x.

When `interval = "Wald"`, an asymptotic $100(1 - \alpha)\%$ confidence interval (where α is equal to `1 - level`) for x_0 is calculated according to Equation (4). The standard error is computed using a First-order Taylor series approximation by calling the `stats` function `numericDeriv` to numerically evaluate the gradient of $\widehat{x}_0$ as a function of y_0 and the fixed-effects $\widehat{\boldsymbol{\beta}}$.

```
invest(transformed.lme, y0 = 500, interval = "Wald")

## estimate    lower    upper       se
##     8.02     4.18    11.85     1.95
```

Alternatively, one can use the very useful `deltaMethod` function from the `car` package to obtain `se`:

```
library(car)  # assuming package car is already installed
params <- c(fixef(transformed.lme), 500)
covmat <- diag(3)   # set up var/cov matrix
covmat[1:2, 1:2] <- vcov(transformed.lme)   # fixed-effects var/cov matrix
covmat[3, 3] <-  17572.35  # VAR[Y_0]
names(params) <- c("b0", "b1", "y0")
deltaMethod(params, g = "(y0 - b0)/b1",  vcov. = covmat)$SE

## [1] 1.95
```

The only drawback here is that `deltaMethod` relies on the `stats` package symbolic differentiation function `D`; hence, $\widehat{x}_0 = f^{-1}\left(y_0; \widehat{\boldsymbol{\beta}}\right)$ has to be obtainable in closed-form.

To obtain the inversion interval (7), we set `interval = "inversion"` (the default) in the call to invest:

```
invest(transformed.lme, y0 = 500, interval = "inversion")

## estimate    lower    upper
##     8.02     4.23    11.92
```

Essentially, `invest` finds the lower and upper inversion confidence limits (7) by solving the equations

$$Q_I - z_{\alpha/2} = 0 \quad \text{and} \quad Q_I - z_{1-\alpha/2} = 0$$

numerically for x using `uniroot`. To use the quantiles from a t distribution instead (see Section 5), we can supply them via the arguments `q1` and `q2`:

```
tvals <- qt(c(0.025, 0.975), df = nrow(bladder) - 1)
invest(transformed.lme, y0 = 500, q1 = tvals[1L], q2 = tvals[2L])

## estimate    lower    upper
##     8.02     4.20    11.95
```

Being able to specify the arguments `q1` and `q2` will also be useful when implementing the bootstrap adjusted inversion interval described in Section **??**

A.2   Using the `lme4` package

Implementation of the parametric bootstrap algorithm in Figure 2 is relatively straight forward using the new `bootMer` function from the well-known R package `lme4` [Bates et al., 2014] in conjunction with the `boot` package.

```
# install.packages("investr", "lme4")
```

Since we will be using the `lme4` package, we need to refit the model using the `lmer` function:

```
library(lme4)
transformed.lmer <- lmer(HD ^ (3 / 2) ~ volume + (0 + 1 | subject) +
                         (0 + volume | subject), data = bladder)
```

Theoretically, the parameter estimates from this model should be the same as those from `transformed.lme`; however, there are likely to be small numerical differences between the two. For this reason, let us re-estimate $\widehat{x}_0$ using `transformed.lmer`. Since `invest` does not work on objects of class `"lmer"`, we have to do things manually:

```
fe <- unname(fixef(transformed.lmer))  # fixed-effects
(x0.est <- (500 - fe[1]) / fe[2])

## [1] 8.02
```

Also, for convenience, we define the following function which estimates $\mathrm{VAR}\left[\mathcal{Y}|x\right] = \sigma_0^2 + x^2\sigma_1^2 + \sigma^2$ for a given value of x:

```
varY <- function(object, x) {
  vc <- as.data.frame(lme4::VarCorr(object))$vcov
  vc[1] + vc[2]*x^2 + vc[3]
}
```

For example, to estimate $\sigma_0^2 = \widehat{\mathrm{VAR}}\left[\mathcal{Y}_0\right]$, we have `varY(transformed.lmer, x = x0.est)`, which gives 1.757×10^4, the same value used in the previous section.

Although we could easily compute all the bootstrap intervals previously discussed in one call to `bootMer` and `boot.ci`, we will discuss and compute each interval separately.

The following snippet of code generates $R = 9999$ bootstrap replicates of $\widehat{x}_0$, Q_W, and Q_I according to the algorithm in Figure 2:

```
pboot <- bootMer(transformed.lmer, nsim = 9, seed = 105, FUN = function(.) {

  # Point estimate
```

```r
var.Y0.boot <- varY(., x = x0.est)    # VAR[Y0]
fe.boot <- unname(fixef(.))    # fixed-effects
if (all(getME(., "y") == bladder$HD ^ (3 / 2))) {
  y0.boot <- 500
} else {
  y0.boot <- rnorm(1, 500, sqrt(var.Y0.boot))
}
x0.boot <- (y0.boot - fe.boot[1]) / fe.boot[2]

# Approximate variance
covmat <- diag(3)
covmat[1L:2L, 1L:2L] <- as.matrix(vcov(.))
covmat[3L, 3L] <-  var.Y0.boot
params <- c("b0" = fe.boot[1],
            "b1" = fe.boot[2],
            "y0" = y0.boot)
dm <- car::deltaMethod(params, g = "(y0 - b0) / b1",  vcov. = covmat)
var.x0.boot <- dm$SE ^ 2

# Approximate predictive pivot
mu0.boot <-  as.numeric(crossprod(fe.boot, c(1, x0.est)))
var.mu0.boot <- t(c(1, x0.est)) %*%
  as.matrix(vcov(.)) %*% c(1, x0.est)
QI.boot <- (y0.boot - mu0.boot) / sqrt(var.Y0.boot + var.mu0.boot)

# Return values
c(x0.boot, var.x0.boot, QI.boot)

})
```

The `bootMer` function returns an object of class `boot` which can then be processed via the `boot` package to obtain the various bootstrap confidence intervals discussed earlier. A basic summary of `pboot` is given by

```r
library(boot)   # load boot package
summary(pboot)

##      R original bootBias bootSE bootMed
## 1 9999     8.02 -0.00219  1.974  8.0186
## 2 9999     1.95 -0.00330  0.256  1.9417
## 3 9999     0.00 -0.00810  1.004  0.0016
```

The estimated standard error and bias of $\widehat{x}_0$, based on `R = 9999` bootstrap replicates, are 1.974 and -0.003, respectively. The original estimate $\widehat{x}_0$ and the median of the bootstrap replicates are also given in the first row of the summary. A graphical summary of the bootstrap simulation is given in Figure 4.

These graphs indicate that the sampling distributions of $\widehat{x}_0$, Q_W, and Q_I are all approximately normal; hence, we would expect the bootstrap confidence intervals to be similar to the asymptotic methods based on the normal distribution discussed in Sections 4-5.
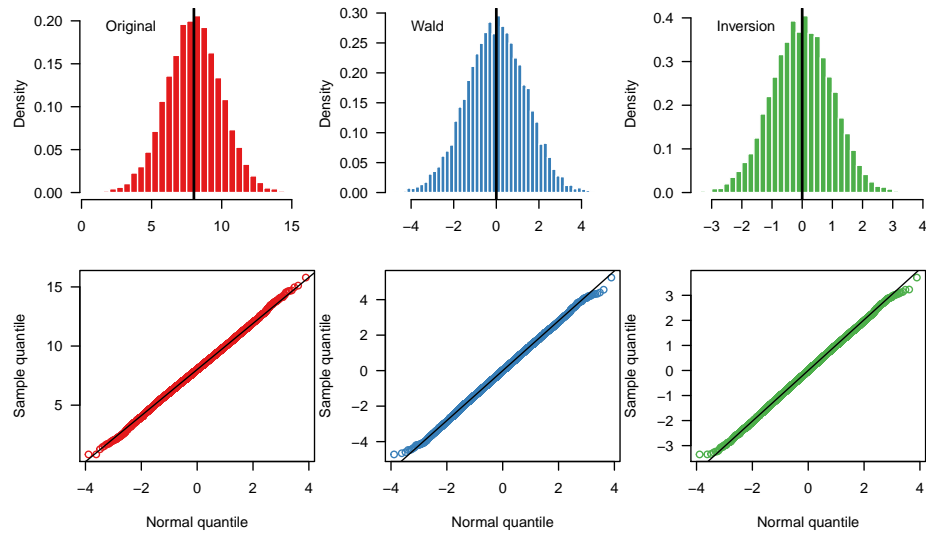

Figure 4: Graphical summary of bootstrap replicates. $R = 9,999$ bootstrap replicates of $\widehat{x}_0$ (left), Q_W (middle), and Q_I (right).

To obtain the percentile and studentized t intervals (see Sections ??-??), we can use the boot package function boot.ci:

```
boot.ci(pboot, type = c("norm", "perc", "stud"))

## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 9999 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = pboot, type = c("norm", "perc", "stud"))
##
## Intervals :
## Level       Normal              Studentized             Percentile
## 95%   ( 4.15, 11.89 )   ( 4.25, 11.95 )   ( 4.05, 11.91 )
## Calculations and Intervals on Original Scale
```

For comparison, we also included the option to compute a bootstrap normal-approximation confidence interval. This interval has the form

$$(\widehat{x}_0 - \texttt{bootBias}) \pm z_{\alpha/2}\texttt{bootSE}$$

16

where `bootBias` and `bootSE` can be found in the first row of `summary(pboot)`. In other words, it is just a Wald-type interval that uses a bias-corrected estimate of x_0, along with a bootstrap estimate of the standard error of $\widehat{x}_0$. While this may be more accurate than the ordinary Wald-based interval (4), it may still not perform well in small sample sizes because of the strict normality assumption. In this example, however, normality does not appear to be an issue.

The bootstrap adjusted inversion interval can be computed as easily as the ordinary inversion interval, except we need to supply `invest` with the estimated quantiles $\widehat{F}_{Q_I}(0.025)$ and $\widehat{F}_{Q_I}(0.975)$:

```
QI.boot <- pboot$t[, 3]   # bootsrap replicates of Q_I
qvals <- quantile(QI.boot, c(0.025, 0.975))   # sample quantiles
invest(transformed.lme, y0 = 500, q1 = qvals[1], q2 = qvals[2])

## estimate    lower    upper
##     8.02     4.28    12.02
```